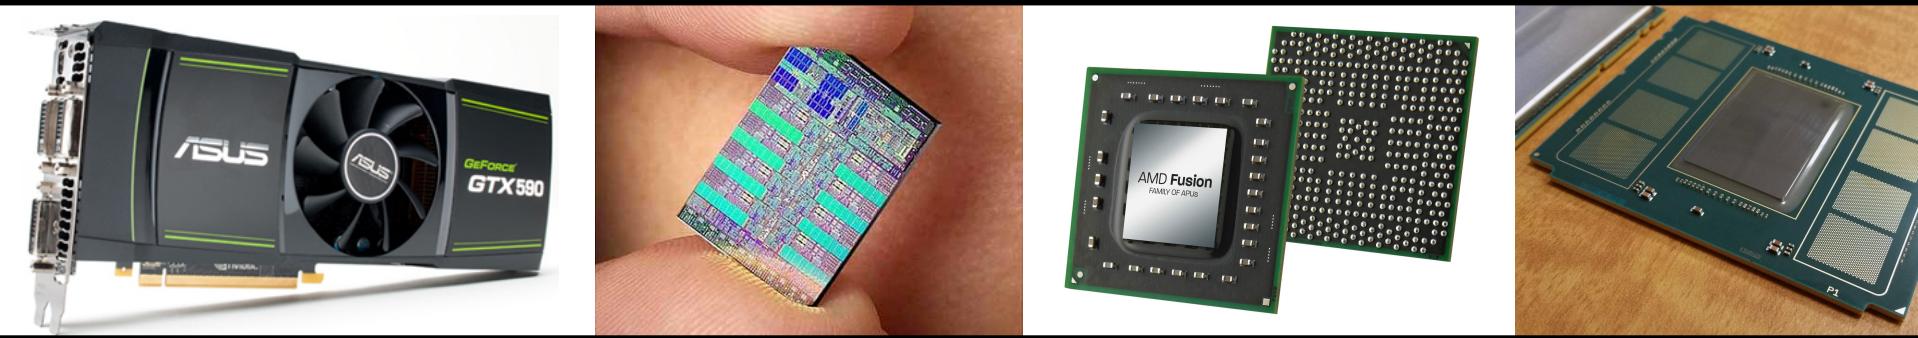


*Exceptional service in the national interest*



## KokkosKernels Overview

S. Rajamanickam,

S. Acer, L. Berger-Vergiat, V. Dang, N. Ellingwood, E. Harvey, B. Kelley, K. Kim, C.R. Trott, J. Wilke

[srajama@sandia.gov](mailto:srajama@sandia.gov)

Center for Computing Research

Sandia National Laboratories, NM

Unclassified Unlimited Release

SAND2020-6470 TR



Sandia National Laboratories is a multimission laboratory managed and operated by National Technology and Engineering Solutions of Sandia, LLC., a wholly owned subsidiary of Honeywell International, Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA-0003525.

# Presentation outline

1. Introduction
2. Dense linear algebra: BLAS capabilities
3. Team level linear algebra: BatchedBLAS
4. Sparse linear algebra, preconditioners and graph kernels
5. Build system using Cmake

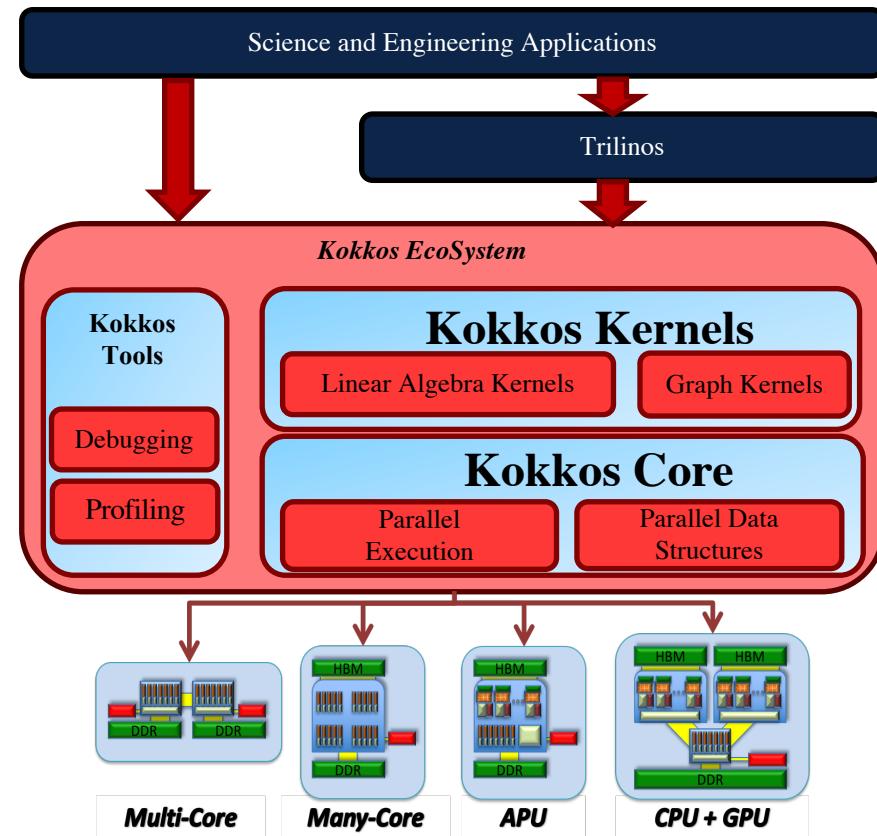
# Presentation outline

1. Introduction
2. Dense linear algebra: BLAS capabilities
3. Team level linear algebra: BatchedBLAS
4. Sparse linear algebra, preconditioners and graph kernels
5. Build system using Cmake

# Goal of the Project

KokkosKernels provides math kernels for dense and sparse linear algebra as well as graph computations. It has multiple aims:

- Portable BLAS, Sparse and Graph kernels
- Generic implementations for various scalar types and data layouts
- Access to major vendor optimized math libraries.
- Expand the scope of BLAS to hierarchical implementations.



# KokkosKernels Tutorial and Exercises



The hands-on exercises that follow demonstrate representative examples from the following categories of capabilities:

- BLAS and Team-based BLAS kernels
- Batched BLAS kernels
- Sparse linear-algebra kernels
- Sparse graph kernels
- Conjugate Gradient examples (Sparse matrix CG and Preconditioned CG)

## Exercise location:

<https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels>

## Preliminary step: run the `run_installlibs_cmake.sh` script

- This will install kokkos, kokkos-kernels, and run CMake to generate a Makefile

The KokkosKernels Wiki will be a useful API reference for the exercises:

<https://github.com/kokkos/kokkos-kernels/wiki/APIReference>

# Presentation outline

1. Introduction
2. Dense linear algebra: BLAS capabilities
3. Team level linear algebra: BatchedBLAS
4. Sparse linear algebra, preconditioners and graph kernels
5. Build system using Cmake

# Capabilities I: BLAS

## BLAS

*BLAS-1 functions are available as multi-vector variants.*

- abs(y,x)	y[i] =  x[i]
- axpy(alpha,x,y)	y[i] += alpha * x[i]
- axpby(alpha,x,beta,y)	y[i] = beta * y + alpha * x[i]
- dot(x,y)	dot = SUM_i ( x[i] * y[i] )
- fill(x,alpha)	x[i] = alpha
- iamax(x)	i = min{i of MAX_i(x[i]) }
- mult(gamma,y,alpha,A,x)	y[i] = gamma * y[i] + alpha * A[i] * x[i]
- nrm1(x)	nrm1 = SUM_i(  x[i]  )
- nrm2(x)	nrm2 = sqrt ( SUM_i(  x[i]  *  x[i]  ) )
- nrm2w(x,w)	nrm2w = sqrt ( SUM_i( ( x[i] / w[i] )^2 ) )
- nrminf(x)	nrminf = MAX_i(  x[i]  )
- reciprocal(r,x)	r[i] = 1 / x[i]
- scal(y,alpha,x)	y[i] = alpha * x[i]
- sum(x)	sum = SUM_i( x[i] )
- update(a,x,b,y,g,z)	y[i] = g * z[i] + b * y[i] + a * x[i]
- gemv(t,alph,A,x,bet,y)	y[i] = bet*y[i] + alph*SUM_j(A[i,j]*x[j])
- gemm(tA,tB,alph,A,B,bet,C)	C[i,j]=bet*C[i,j]+alph*SUM_k(A[i,k]*B[k,j])
- trmm(s,uplo,t,d,alpha,A,B)	B = alpha*op(A)*B or alpha*B*op(A)
- trsm(s,uplo,t,d,alpha,A,B)	X = op(A)\alpha*B or alpha*B/op(A)

# Capabilities II: LAPACK and TeamBLAS



## BLAS/LAPACK

- GESV
- TRTRI

## TEAM

- abs
- axpby
- dot
- mult
- norm2
- scal
- update

# Capabilities I: BLAS

## BLAS

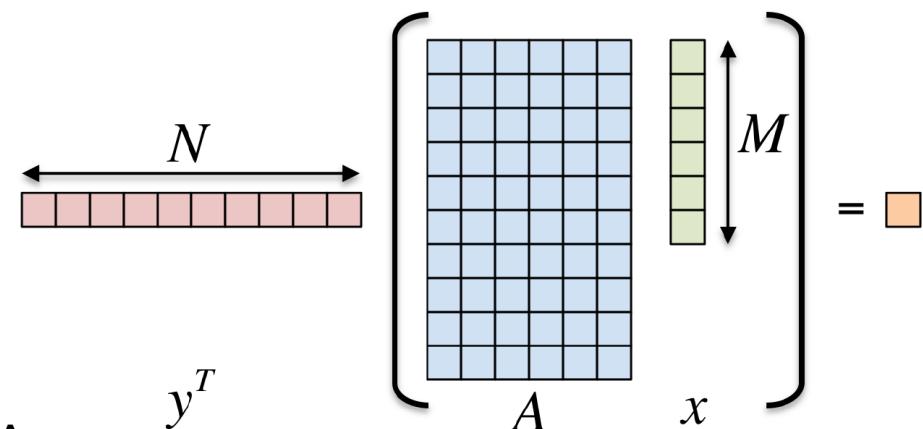
*BLAS-1 functions are available as multi-vector variants.*

- abs(y,x)	y[i] =  x[i]
- axpy(alpha,x,y)	y[i] += alpha * x[i]
- axpby(alpha,x,beta,y)	y[i] = beta * y + alpha * x[i]
- dot(x,y)	dot = SUM_i ( x[i] * y[i] )
- fill(x,alpha)	x[i] = alpha
- iamax(x)	i = min{i of MAX_i(x[i]) }
- mult(gamma,y,alpha,A,x)	y[i] = gamma * y[i] + alpha * A[i] * x[i]
- nrm1(x)	nrm1 = SUM_i(  x[i]  )
- nrm2(x)	nrm2 = sqrt ( SUM_i(  x[i]  *  x[i]  ) )
- nrm2w(x,w)	nrm2w = sqrt ( SUM_i( ( x[i] / w[i] )^2 ) )
- nrminf(x)	nrminf = MAX_i(  x[i]  )
- reciprocal(r,x)	r[i] = 1 / x[i]
- scal(y,alpha,x)	y[i] = alpha * x[i]
- sum(x)	sum = SUM_i( x[i] )
- update(a,x,b,y,g,z)	y[i] = g * z[i] + b * y[i] + a * x[i]
- gemv(t,alph,A,x,bet,y)	y[i] = bet*y[i] + alph*SUM_j(A[i,j]*x[j])
- gemm(tA,tB,alph,A,B,bet,C)	C[i,j]=bet*C[i,j]+alph*SUM_k(A[i,k]*B[k,j])
- trmm(s,uplo,t,d,alpha,A,B)	B = alpha*op(A)*B or alpha*B*op(A)
- trsm(s,uplo,t,d,alpha,A,B)	X = op(A)\alpha*B or alpha*B/op(A)

# KokkosKernels Interface: Inner Product Exercise

## BLAS and Team-based BLAS interface

- Goal: Re-implement the inner product exercise using
  1. KokkosKernels BLAS functions
  2. KokkosKernels team-based BLAS function
- Details:  $\langle \mathbf{y}, \mathbf{A}^* \mathbf{x} \rangle$ 
  - $\mathbf{y}$  is  $N \times 1$ ,  $\mathbf{A}$  is  $N \times M$ ,  $\mathbf{x}$  is  $M \times 1$
  - Look for comments labeled with “EXERCISE”
  - Compile and run on OpenMP, CUDA backends
  - Compare performance of these two implementations
  - Try LayoutLeft and LayoutRight



# KokkosKernels Interface: Inner Product Exercise

- Exercise location:

<https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/InnerProduct/Begin>

- Details:

1. KokkosKernels BLAS functions

- Convert from hierarchical parallel execution to using BLAS functions
- **tmp = A\*x** (tmp for **gemv** results)
- **result = <y,tmp>**

2. KokkosKernels team-based BLAS

- Same as hierarchical parallel execution
- Call team-based **dot** within each team to perform **<A[teamId,:],x>**

## Hint: KokkosKernels Functions

```
result = KokkosBlas::dot(x,y)
```

performs `result = SUM_i(y[i]*x[i])`

```
KokkosBlas::gemv("N",alpha,A,x,beta,y)
```

performs matrix-vector multiplication

```
y[i] = beta*y[i] +  
alpha*SUM_j(A[i,j]*x[j])
```

```
KokkosBlas::Experimental::dot(teamId,x,y)
```

performs `result = SUM_i(y[i]*x[i])` within each thread team

# Presentation outline

1. Introduction
2. Dense linear algebra: BLAS capabilities
3. Team level linear algebra: BatchedBLAS
4. Sparse linear algebra, preconditioners and graph kernels
5. Build system using Cmake

# Capabilities III: Batched BLAS/LAPACK



## Batched BLAS

```
parallel_for(TeamPolicy(N), KOKKOS_LAMBDA(member_type &member) {  
    parallel_for(TeamThreadRange(member, nRows), [&](int &i) {  
        parallel_for(ThreadVectorRange(member, nCols), [&](cint &j) {  
            for (int k=0;k<len;++k)  
                C(i,j) += A(i,k) * B(k,j);  
        });});});});});
```

← TeamVectorGemm  
← SerialDot

- Parallel batched operations are composed of using parallel-for with hierarchical interface of BLAS/LAPACK.
- Provides hierarchical interface matching to Kokkos hierarchical parallelism.
  - Serial – a single thread can invoke this layer of functions
  - Team – a single level of nested parallelism is used
  - TeamVector
    - Two nested parallel-for's are used TeamThread, ThreadVector, or
    - A single flattened nested parallel-for is used with TeamVectorRange.
- Function list
  - [S] Givens rotation, [S,TV] Householder transformation, [S] Hessenberg transformation, [S,T] LU, [S,T] Gemm, [S,T] Gemv, [S,T] Trsm, [S,T] Trsv, [S] Trmm, [S] Trtri, etc.
  - Implement new functions upon requests.

# Capabilities III: Batched BLAS/LAPACK



## Batched BLAS

```
parallel_for(TeamPolicy(N), KOKKOS_LAMBDA(member_type &member) {  
    parallel_for(TeamThreadRange(member, nRows), [&](int &i) {  
        parallel_for(ThreadVectorRange(member, nCols), [&](cint &j) {  
            for (int k=0;k<len;++k)  
                C(i,j) += A(i,k) * B(k,j);  
        });});});});});
```

← TeamVectorGemm  
← SerialDot

- Parallel batched operations are composed of using parallel-for with hierarchical interface of BLAS/LAPACK.
- Provides hierarchical interface matching to Kokkos hierarchical parallelism.
  - Serial – a single thread can invoke this layer of functions
  - Team – a single level of nested parallelism is used
  - TeamVector
    - Two nested parallel-for's are used TeamThread, ThreadVector, or
    - A single flattened nested parallel-for is used with TeamVectorRange.
- Function list
  - [S] Givens rotation, [S,TV] Householder transformation, [S] Hessenberg transformation, [S,T] LU, [S,T] Gemm, [S,T] Gemv, [S,T] Trsm, [S,T] Trsv, [S] Trmm, [S] Trtri, etc.
  - Implement new functions upon requests.

# Team-level Batched BLAS: Gemm

- Objective
  - Learn how to use the Team-level interface for Gemm.
- Exercise location:

<https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/TeamGemm/Begin>
- Details:
  1. Perform general matrix multiplies in parallel on a batch of matrices using the team-level Gemm function.
  2. Tune the team size on Cuda (via  $-A.m$ ) and see how it impacts the performance.

# KokkosBatched BLAS/LAPACK

## Construction & Apply BlockJacobi

- Objective
  - Compose a batched LU factorization on diagonal blocks and compute the inverse of the blocks.
  - Compare different parallelization strategies: a sequence of parallel\_for calls vs a single parallel\_for interface.
- Exercise location:

<https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/BlockJacobi/Begin>
- Details:
  1. Simulate the standard batch functions running a sequence of parallel\_for with LU, Lower Trsm, Upper Trsm to invert diagonal blocks.
  2. Compose a batched block matrix inversion using the so-called batched BLAS/LAPACK interface and a parallel\_for with functor-level functions.
  3. On GPUs,
    1. Test the code with different team size with the run-different-teamsize.sh script.
    2. Run nvprof and see the difference (run-nvprof.sh script) and tune the team size for the problem size.

# Presentation outline

1. Introduction
2. Dense linear algebra: BLAS capabilities
3. Team level linear algebra: BatchedBLAS
4. Sparse linear algebra, preconditioners and graph kernels
5. Build system using Cmake

# Capabilities IV: Sparse Kernels

## Sparse

- CSR-Sparse Matrix Class providing fundamental capabilities
- SPMV: Sparse Matrix Vector Multiply
- SPADD: Sparse Matrix Matrix Addition
- SPGEMM: Sparse Matrix Matrix Multiply; separate symbolic and numeric phase
- SPGEMM\_Jacobi: Fused Sparse Matrix Matrix Multiply + Jacobi
- GS: Gauss-Seidel Method using graph coloring: symbolic, numeric, solve phases]
- SPILU(K): Incomplete LU Factorization
- SPTRSV: Sparse triangular solve; super-nodal variant

## Graph

- Distance-1 and Distance-2 graph coloring
- Triangle enumeration for graph analytics
  - Using SpGEMM + Visitor Pattern: can be used to represent large problems

# SPARSE: CrsMatrix

- Example location:  
[https://github.com/kokkos/kokkos-kernels/blob/develop/example/wiki/sparse/KokkosSparse\\_wiki\\_crsmatrix.cpp](https://github.com/kokkos/kokkos-kernels/blob/develop/example/wiki/sparse/KokkosSparse_wiki_crsmatrix.cpp)
- Container for compressed row sparse matrix (CrsMatrix)
  - API available in our wiki: <https://github.com/kokkos/kokkos-kernels/wiki/CrsMatrix>
  - Currently no deep\_copy constructor or operator
  - Row pointers is const, you cannot re-order in place or increase storage space
- Interfaced with the following kernels:
  - SpMV
  - SpADD
  - SpGEMM

# SPARSE: SPMV

- Exercise location:  
[https://github.com/kokkos/kokkos-kernels/blob/develop/example/wiki/sparse/KokkosSparse\\_wiki\\_spmv.cpp](https://github.com/kokkos/kokkos-kernels/blob/develop/example/wiki/sparse/KokkosSparse_wiki_spmv.cpp)
- Usage:
  - `KokkosSparse::spmv(mode, alpha, A, x, beta, y);`
- Performs:
  - $y = \beta * y + \alpha * \text{op}(A) * x$
- Parameters:
  - mode: input “N” (non-transpose), “T” (transpose), “C” (conjugate-transpose)
  - alpha, beta: input scalars
  - A: input sparse matrix (`KokkosSparse::CrsMatrix`) of size  $n \times m$
  - x: input vector (rank-1 or rank-2 `Kokkos::View`) of length  $m$
  - y: input/output vector (rank-1 or rank-2 `Kokkos::View`) of length  $n$

# SPARSE: SPADD

- Exercise location:
  - [https://github.com/kokkos/kokkos-kernels/blob/develop/example/wiki/sparse/KokkosSparse\\_wiki\\_spadd.cpp](https://github.com/kokkos/kokkos-kernels/blob/develop/example/wiki/sparse/KokkosSparse_wiki_spadd.cpp)
- Usage:
  - KokkosSparse::spadd\_symbolic(handle, A, B, C);
  - KokkosSparse::spadd\_numeric(handle, alpha, A, beta, B, C);
- Performs:
  - $C = \text{beta} * C + \text{alpha} * (A+B)$
- Parameters:
  - $\text{alpha}, \text{beta}$ : input scalars
  - $A, B$ : input sparse matrices (KokkosSparse::CrsMatrix) of size  $n \times m$
  - $C$ : input/output sparse matrix (KokkosSparse::CrsMatrix) of size  $n \times m$

# SPARSE: SpGEMM

- Exercise location:  
<https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/SpGEMM/Begin>
- **Usage:**
  - KokkosSparse::spgemm\_symbolic(handle, A, isTransposed, B, isTransposed, C);
  - KokkosSparse::spgemm\_numeric(handle, A, isTransposed, B, isTransposed, C);
- **Performs:**
  - $C = \text{op}(A) * \text{op}(B)$
- **Notes:**
  - Create an SpGEMM handle:  
`handle.create_spgemm_handle(spgemm_algorithm);`
  - Create an empty matrix for output: KokkosSparse::CrsMatrix<...> C;
  - Use the same handle for both symbolic and numeric calls

# SPARSE: Gauss-Seidel

Given: **A**, a CrsMatrix that is numRows x numCols

**x**: unknown vector. **b**: right-hand side vector.

First, create KokkosKernelsHandle and call `create_gs_handle()`.

Structural (symbolic) setup.

- `KokkosSparse::Experimental::gauss_seidel_symbolic(&handle, numRows, numCols, A.graph.row_map, A.graph.entries, graphIsSymmetric);`

Numeric setup: anything depending on matrix values.

- `KokkosSparse::Experimental::gauss_seidel_numeric(&handle, numRows, numCols, A.graph.row_map, A.graph.entries, A.values, graphIsSymmetric);`

Forward apply (symmetric and backward have same interface):

- `KokkosSparse::Experimental::forward_sweep_gauss_seidel_apply(&handle, numRows, numCols, A.graph.row_map, A.graph.entries, A.values, x, b, initZeroX, updateCachedB, omega, numSweeps);`

# SPARSE: Gauss-Seidel Exercise

- Why both numRows and numCols?
  - For distributed Tpetra::CrsMatrix, numRows is number of locally owned rows, and numCols is locally owned + remote columns (column map). Matrix should still be globally square.
- initZeroX: whether to zero out x before apply. Only set to true on first apply on a given system.
- updateCachedB: if first apply with this b, or b has changed for any reason since last apply.

**Exercise:** <https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/GaussSeidel/Begin>

- First, fill in the code after // EXERCISE comments. Check that the solve converges in ~7 iterations.
- Then: try other algorithms when creating the GS sub-handle:

```
handle.create_gs_handle(KokkosSparse::GS_CLUSTER, 10);
```

or:

```
handle.create_gs_handle(KokkosSparse::GS_TWOSTAGE);
```
- In apply, try increasing number of sweeps from 1, or omega (damping factor) to something other than 1.0
- Try using symmetric\_gauss\_seidel\_apply instead of forward

# GRAPH: Coloring

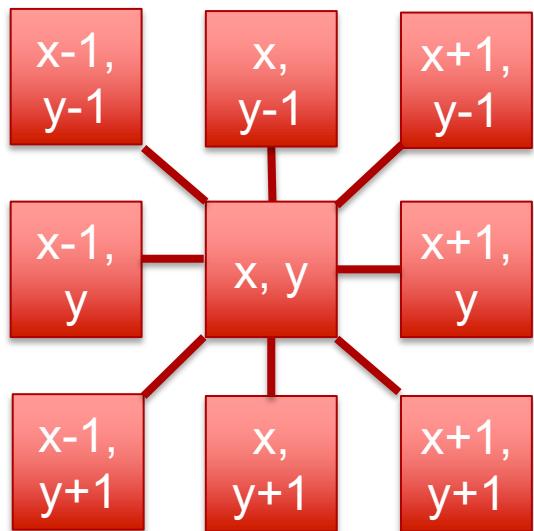
Given: **rowmap**, **entries** (forming a sparse CRS graph)

Compute distance-1 coloring: no pair of adjacent vertices can have same color. Example usage of bitwise vertex-based (VBBIT) algorithm:

```
// Set up
KernelHandle handle;
handle.create_graph_coloring_handle(KokkosGraph::COLORING_VB);
// Compute the coloring
KokkosGraph::Experimental::graph_color(&handle, numVertices,
numVertices, rowmap, entries);
// Get the colors array and number of colors used
auto colorHandle = handle.get_graph_coloring_handle();
auto colors = colorHandle->get_vertex_colors();
auto numColors = colorHandle->get_num_colors();
// Clean up
handle.destroy_graph_coloring_handle();
```

# GRAPH: Coloring Exercise

- **Exercise:** <https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/GraphColoring/Begin>
- Creates graph from 2D rectangular grid and a 9-point stencil



- Distance-2 coloring: vertices that are within 1 or 2 edges of each other can't have same color.
- Exercise runs both D-1 and D-2 coloring, and prints out the colors in a grid.
- (If building with OpenMP) Try changing the number of threads and see how the number of colors changes:
- `export OMP_NUM_THREADS=<n>`

- **Some D1 algorithms:**  
`COLORING_VB`, `COLORING_SERIAL`, `COLORING_EB`
- **Some D2 algorithms:**  
`COLORING_D2_VB`, `COLORING_D2_SERIAL`, `COLORING_D2_NB_BIT`

# SPARSE: SpILUK



- Goal: get familiar with Kokkos Kernels interfaces of ILU(k) for CRS matrices
  - Exercise location:

<https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/SplLUK/Begin>

- Kokkos Kernels interfaces:

# SPARSE: SpILUK

- Exercise Instructions:
  - ❑ Generate a simple Laplacian matrix on a cartesian grid as a KokkosSparse::CrsMatrix
  - ❑ Create an SpILUK handle
  - ❑ Call `spiluk_symbolic(...)` in the symbolic phase to construct the nonzero patterns of L and U, and to do level scheduling
  - ❑ Call `spiluk_numeric(...)` in the numeric phase to perform an ILU factorization based on the nonzero patterns found in the symbolic phase
  - ❑ As fill-level  $k$  changes, observe the changes of
    - ILU( $k$ ) fill-factor, and
    - row-sum difference ( $A^* \text{ones} - L^* U^* \text{ones}$ )

# KokkosKernels Interface: Conjugate Gradient Solver



- Goal: implement conjugate gradient solver for square, symmetric, positive-definite sparse matrix
- Details:  $\mathbf{A}^* \mathbf{x} = \mathbf{b}$ 
  - $\mathbf{b}$  is  $N \times 1$
  - $\mathbf{A}$  is  $N \times N$  symmetric, positive-definite sparse matrix
  - $\mathbf{x}$  is  $N \times 1$
  - Look for comments labeled with “EXERCISE”
  - Use KokkosKernels BLAS and KokkosKernels Sparse BLAS

- Algorithm:

$$\mathbf{r}_0 = \mathbf{b} - \mathbf{A} * \mathbf{x}_0$$

$$\mathbf{p}_0 = \mathbf{r}_0$$

$$k = 0$$

**while**  $\|\mathbf{r}_k\| > \varepsilon$  and  $k < N$

$$\alpha = \frac{\mathbf{r}_k^T * \mathbf{r}_k}{\mathbf{p}_k^T * \mathbf{A} * \mathbf{p}_k}$$

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha * \mathbf{p}_k$$

$$\mathbf{r}_{k+1} = \mathbf{r}_k - \alpha * \mathbf{A} * \mathbf{p}_k$$

$$\beta = \frac{\mathbf{r}_{k+1}^T * \mathbf{r}_{k+1}}{\mathbf{r}_k^T * \mathbf{r}_k}$$

$$\mathbf{p}_{k+1} = \mathbf{r}_{k+1} + \beta * \mathbf{p}_k$$

$$k = k + 1$$

# KokkosKernels Interface: Conjugate Gradient Solver

- Exercise location:

<https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/CGSolve/Begin>

- Details:

- Sparse matrix generation is provided
- Compile and run on OpenMP, CUDA backends
- Vary problem size: -N #
- Compare performance of CPU vs GPU

## Hint: KokkosKernels Functions

**result = KokkosBlas::dot(x,y)**

performs  $\text{result} = \text{SUM}_i (y[i]^*x[i])$

**KokkosBlas::axpy(alpha,x,y)**

performs  $y[i] = y[i] + \alpha*x[i]$

**KokkosBlas::axpby(alpha,x,beta,y)**

performs  $y[i] = \beta*y[i] + \alpha*x[i]$

**KokkosSparse::spmv("N",alpha,A,x,beta,y)**

performs sparse matrix-vector multiplication

$y[i] = \beta*y[i] + \alpha*\text{SUM}_j (A[i,j]^*x[j])$

# ILU(k) Preconditioned Conjugate Gradient Solver with Kokkos Kernels



- Goal: implement an ILU(k) preconditioned conjugate gradient solver for square, symmetric, positive-definite sparse matrix using Kokkos Kernels SpILUK and SpTRSV

- Exercise location:

[https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/CGSolve\\_SpILUKprecond /Begin](https://github.com/kokkos/kokkos-tutorials/tree/master/Intro-Full/Exercises/kokkoskernels/CGSolve_SpILUKprecond /Begin)

- Kokkos Kernels interfaces:

# ILU(k) Preconditioned Conjugate Gradient Solver with Kokkos Kernels

- Exercise Instructions:

- Generate a simple Laplacian matrix on a cartesian grid as a KokkosSparse::CrsMatrix
- Create two SpTRSV handles (L and U)
- Call `sptrsv_symbolic(...)` to do level scheduling
- Call `sptrsv_solve(...)` to apply the preconditioner during the CGSolve

1. `tmp = L\r`
2. `z = U\ (tmp)`

- Observe the convergence behaviors:
  - without preconditioner
  - with preconditioner (as ILU(k) fill-level changes)

- Algorithm:

$$\mathbf{r}_0 = \mathbf{b} - \mathbf{A} * \mathbf{x}_0$$

$$\mathbf{z}_0 = \mathbf{M}^{-1} * \mathbf{r}_0$$

$$\mathbf{p}_0 = \mathbf{z}_0$$

$$k = 0$$

**while**  $\|\mathbf{r}_k\| > \varepsilon$  and  $k < N$

$$\alpha = \frac{\mathbf{r}_k^T * \mathbf{z}_k}{\mathbf{p}_k^T * \mathbf{A} * \mathbf{p}_k}$$

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha * \mathbf{p}_k$$

$$\mathbf{r}_{k+1} = \mathbf{r}_k - \alpha * \mathbf{A} * \mathbf{p}_k$$

$$\mathbf{z}_{k+1} = \mathbf{M}^{-1} * \mathbf{r}_{k+1}$$

$$\beta = \frac{\mathbf{r}_{k+1}^T * \mathbf{z}_{k+1}}{\mathbf{r}_k^T * \mathbf{z}_k}$$

$$\mathbf{p}_{k+1} = \mathbf{z}_{k+1} + \beta * \mathbf{p}_k$$

$$k = k + 1$$

# Presentation outline

1. Introduction
2. Dense linear algebra: BLAS capabilities
3. Team level linear algebra: BatchedBLAS
4. Sparse linear algebra, preconditioners and graph kernels
5. Build system using Cmake

# Building Kokkos Kernels with CMake

- `cmake ${KERNELS_SOURCE} -D{OPTION}:BOOL=ON ...`  
Specify config variables on cmd line as `-DOPT=Value`
- To get a list of options in TUI, use `ccmake`:
  - `ccmake ${KERNELS_SOURCE}`  
`-DCMAKE_CXX_COMPILER={}`  
`-DKokkos_ROOT=${PATH_TO_KOKKOS_INTALL}`

```

KokkosKernels_INST_OFFSET_SIZE          ON
KokkosKernels_INST_ORDINAL_INT         ON
KokkosKernels_INST_ORDINAL_INT         OFF
KokkosKernels_LAPACKE_ROOT             OFF
KokkosKernels_LAPACK_ROOT              ON
KokkosKernels_LINALG_OPT_LEVEL        ON
KokkosKernels_MAGMA_ROOT               ON
KokkosKernels_METIS_ROOT               OFF
KokkosKernels_MKL_ROOT                ON
KokkosKernels_NO_DEFAULT_CUDA_         OFF
KokkosKernels_SUPERLU_ROOT             ON
KokkosKernels_TEST_ETI_ONLY           ON
Kokkos_DIR                           /Users/jjwilke/Programs/install/kokkos/develop/openmp/lib/cmake/Kokkos
LAPACKE_INCLUDE_DIRS                  LAPACKE_LIBRARIES
LAPACKE_LIBRARY_DIRS                 LAPACK_INCLUDE_DIRS
LAPACK_INCLUDE_DIRS                  LAPACK_LIBRARIES
LAPACK_LIBRARY_DIRS                 METIS_INCLUDE_DIRS
METIS_INCLUDE_DIRS                   METIS_LIBRARIES

```

```

METIS_LIBRARIES: Optional override for the libraries that comprise TPL METIS. Default: None. Default common
Press [enter] to edit option Press [d] to delete an entry
Press [c] to configure
Press [h] for help           Press [q] to quit without generating
Press [t] to toggle advanced mode (Currently Off)

```

- Scroll over option and hit ‘h’ to see option details
- Hit ‘enter’ to change option value

# Most Important Build Options

- -DKokkos\_ROOT=\${PATH\_TO\_KOKKOS}
  - The underlying Kokkos. Many options (CUDA vs OpenMP, e.g.) automatically derived from Kokkos installation
- Template instantiation ahead of time to save compile-time
  - Scalars: e.g. -D KokkosKernels\_INST\_DOUBLE:BOOL=ON
  - Offsets: e.g. -D KokkosKernels\_INST\_OFFSET\_SIZE\_T:BOOL=ON
  - Mem/Exec Spaces: e.g.
    - D KokkosKernels\_INST\_EXECSpace\_OPENMP:BOOL=ON
- Third-party library (TPLs)
  - Enable: e.g. -D KokkosKernels\_ENABLE\_TPL\_BLAS:BOOL=ON
    - Kernels will look for BLAS in default locations
  - Customize paths:
    - e.g. -D BLAS\_LIBRARIES="..."
    - e.g. -D BLAS\_LIBRARY\_DIRS="..."

# A basic CMake project file creates a target and links to Kokkos

Target linking in CMake automatically adds all include/link options.  
Only need to link to KokkosKernels (Kokkos automatically comes with)

```
cmake_minimum_required(VERSION 3.12) #need >= 3.12
project(myProject CXX) #kokkos is always a C++ project
```

```
#must find installed kokkos kernels
find_package(KokkosKernels REQUIRED)
```

```
add_executable(myExe source.cpp) #add my program
```

```
#link to kokkos
target_link_libraries(myExe PRIVATE Kokkos::kokkoskernels)
```



Sandia  
National  
Laboratories

*Exceptional service in the national interest*

<http://www.github.com/kokkos>